


# How Does Gradient Descent Work?



Jeremy Cohen · Peking University · Apr 10, 2025

# This talk

- Neural networks are trained using optimization algorithms
- Yet, optimization theory is not used in deep learning. Why?
- Thesis of this talk:
  1. Existing optimization theory does not apply in deep learning ...
  2. ... but a different kind of theory is possible.
- Goal: convince you to help build the theory of optimization in deep learning

# Gradient descent


- The simplest optimizer is deterministic gradient descent (GD):

$$w_{t+1} = w_t - \eta \nabla L(w_t)$$

- Existing theory can't explain the convergence of even this algorithm
- We must understand GD before we can understand more complex methods

# Warm-up: quadratic objective functions

- On quadratics, GD oscillates if the *curvature* (2nd derivative) is too high
- Consider a 1d quadratic function  $L(x) = \frac{1}{2}Sx^2$ , with curvature  $L''(x) = S$




# Warm-up: quadratic objective functions

- For a quadratic in *multiple* dimensions, curvature is quantified by Hessian
- GD oscillates along Hessian eigenvectors with eigenvalues greater than  $2/\eta$

# Warm-up: quadratic objective functions

- For a quadratic in *multiple* dimensions, curvature is quantified by Hessian
- GD oscillates along Hessian eigenvectors with eigenvalues greater than  $2/\eta$




# What about deep learning?

- For DL objectives, can take quadratic Taylor approximation around any  $w$
- Dynamics of GD on this quadratic depend on the top eigenvalue of the Hessian  $H(w)$ , i.e. the *sharpness*  $S(w) := \lambda_1(H(w))$
- If sharpness  $S(w) > 2/\eta$ , GD would diverge on the quadratic Taylor approximation
- This suggests that GD doesn't function properly if sharpness  $S(w) > 2/\eta$


# Gradient descent in deep learning

- Why does gradient descent converge in deep learning?
- Natural idea: sharpness  $S(w)$  remains below  $2/\eta$  throughout training
  - i.e. GD stays inside the “stable region”  $\{w : S(w) \leq 2/\eta\}$



# Gradient descent in deep learning

- Why does gradient descent converge in deep learning?
- Natural idea: sharpness  $S(w)$  remains below  $2/\eta$  throughout training
  - i.e. GD stays inside the “stable region”  $\{w : S(w) \leq 2/\eta\}$





- This is the picture suggested by traditional optimization theory (“L-smoothness”)

# Deep learning reality

- Train neural network using GD with  $\eta = 0.02$  (ViT on CIFAR-10):



# Deep learning reality

- Train neural network using GD with  $\eta = 0.02$  (ViT on CIFAR-10):




Quadratic Taylor approximation predicts growing  
oscillations along top Hessian eigenvector

# What happens next?






What happens next?


This quantity is  
predicted to oscillate





# What happens next?



Mystery Hint: ~~why does the sharpness drop?~~





# Full gradient descent trajectory




# What if we train at a different learning rate?


- Train same network with smaller learning rate  $\eta = 0.01$  (orange):



# Expectation vs. reality




Expectation




Reality

Gradient descent trains at the **edge of stability**

# This behavior is generic across DL settings



# This behavior is generic across DL settings



- This is not a weird edge case, it's the **typical** behavior of GD in DL

# Same phenomenon

Wu, Ma, E. *How SGD Selects the Global Minima in Over-parameterized Learning: A Dynamical Stability Perspective*. NeurIPS '18.

$\eta$	0.01	0.05	0.1	0.5	1	5
FashionMNIST	$53.5 \pm 4.3$	$39.3 \pm 0.5$	$19.6 \pm 0.15$	$3.9 \pm 0.0$	$1.9 \pm 0.0$	$0.4 \pm 0.0$
CIFAR10	$198.9 \pm 0.6$	$39.8 \pm 0.2$	$19.8 \pm 0.1$	$3.6 \pm 0.4$	-	-
prediction $2/\eta$	200	40	20	4	2	0.4

Observation: sharpness at end of training is  $\approx 2/\eta$

# What's going on?

Cohen, Kaur, Li, Kolter, Talwalkar. *Gradient descent on neural networks typically occurs at the edge of stability*. ICLR '21.

Why does gradient descent work in deep learning?

# The answer




Damian\*, Nichani\*, Lee. *Self-stabilization: the implicit bias of gradient descent at the edge of stability.* ICLR '23.

- To understand dynamics of GD, need to Taylor expand to *third-order*.
- This expansion reveals the key ingredient missing from traditional theory:

Oscillations along the top Hessian eigenvector automatically reduce the top Hessian eigenvalue.

# Informal sketch



cartoon of weight-space dynamics

Suppose that GD is oscillating along the top Hessian eigenvector  $u$


How does the gradient  $\nabla L$  at

$$w = w^* + xu$$

relate to the gradient at  $w^*$ ?

# Informal sketch

current iterate  $w$




By Taylor expansion around  $w^*$ :

$$\nabla L(w^* + xu) =$$

gradient at  $w$

# Informal sketch

current iterate  $w$



By Taylor expansion around  $w^*$ :

$$\nabla L(w^* + xu) = \nabla L(w^*) + O(x)$$


gradient at  $w$       gradient at  $w^*$

Since  $u$  is a Hessian eigenvector

# Informal sketch

$$H(w^*) u = S(w^*) u$$

current iterate  $w$



By Taylor expansion around  $w^*$ :

$$\nabla L(w^* + xu) = \underbrace{\nabla L(w^*)}_{\text{gradient at } w} + \underbrace{H(w^*)[xu]}_{= S(w^*) x u} + O(x^2)$$

gradient at  $w^*$       oscillation

- This term sends a negative gradient step computed at  $w^* + xu$  towards the  $-u$  direction.
- This term is causing us to oscillate
- The “magic” comes from the *next* term in the Taylor expansion...

# Informal sketch

- The next term in the Taylor expansion is:

$$\nabla L(w^* + xu) = \underbrace{\nabla L(w^*)}_{\text{gradient at } w} + \underbrace{H(w^*)[xu]}_{\text{gradient at } w^*} + \underbrace{\frac{1}{2} x^2 \nabla_{w^*} [u^T H(w^*) u]}_{\text{oscillation}} + O(x^3)$$

*curvature in  $u$  direction =  $S(w^*)$*

*gradient of curvature in  $u$  direction =  $\nabla S(w^*)$*

# Informal sketch

- The next term in the Taylor expansion is:

$$\nabla L(w^* + xu) = \underset{\text{gradient at } w}{\nabla L(w^*)} + \underset{\text{gradient at } w^*}{H(w^*)[xu]} + \frac{1}{2} x^2 \underset{\text{oscillation}}{\nabla S(w^*)} + O(x^3)$$




gradient of sharpness

- Thus, a negative gradient step computed at  $w^* + xu$  automatically takes a negative gradient step *on the sharpness* with step size  $\frac{1}{2}\eta x^2$ .
- i.e. **oscillations automatically trigger reduction of sharpness**
  - the size of this effect is proportional to the squared magnitude of oscillation
- This is the crucial ingredient missing from the traditional theory.

# Let's revisit the behavior of GD

- When GD exits the stable region:
  - it oscillates along the top Hessian eigenvector (as expected)
  - these oscillations implicitly perform gradient descent *on the sharpness* (top Hessian eigenvalue)
  - this reduces sharpness, thereby steering GD back into the stable region

# Let's revisit the behavior of GD




when  $x$  grows large,  
the effect becomes  
non-negligible


when  $x$  is small,  
the effect is small

# Cause and effect

Traditional theory



# Cause and effect






- Traditional optimization theory fails to capture the **causal structure** of the optimization process
- GD doesn't converge because the curvature is “*a priori*” small — it converges due to an **automatic negative feedback mechanism** that keeps the curvature small.

# How can we analyze gradient descent?

- Unfortunately, EOS dynamics are challenging to analyze in fine-grained detail
- Need to track the mutual interactions between oscillations and curvature
- There are frequently *multiple* unstable eigenvalues => chaotic dynamics

# How can we analyze gradient descent?

- Unfortunately, EOS dynamics are challenging to analyze in fine-grained detail
- Need to track the mutual interactions between oscillations and curvature
- There are frequently *multiple* unstable eigenvalues => chaotic dynamics



# How can we analyze gradient descent?

Cohen\*, Damian\*, Talwalkar, Kolter, Lee. *Understanding Optimization in Deep Learning with Central Flows*. ICLR '25.



Alex Damian


- We argue that the exact oscillatory GD trajectory doesn't matter
- Rather, what matters is the *macroscopic* path that GD takes
- This macroscopic path turns out to be much easier to understand
  - We only need to understand the oscillations in a *statistical* sense

# What path does gradient descent take?

- The standard continuous-time approximation to GD is *gradient flow*:

$$\frac{dw}{dt} = -\eta \nabla L(w)$$

- GD follows gradient flow *before* EOS, but then takes a different path




# What path does gradient descent take?

- The standard continuous-time approximation to GD is *gradient flow*:

$$\frac{dw}{dt} = -\eta \nabla L(w)$$

- GD follows gradient flow *before* EOS, but then takes a different path
- Our *central flow* matches the trajectory of GD even at EOS.




# What path does gradient descent take?

- The standard continuous-time approximation to GD is *gradient flow*:

$$\frac{dw}{dt} = -\eta \nabla L(w)$$


- GD follows gradient flow *before* EOS, but then takes a different path

- Our *central flow* matches the trajectory of GD even at EOS.



# Central flow

- The central flow models the *time-averaged* (i.e. smoothed) GD trajectory



# Deriving the central flow

- We derive the central flow using informal mathematical reasoning, and we show empirically that this flow matches the real GD trajectory
- In particular:
  - We suppose that the time-averaged trajectory can be described by a flow
  - We argue that only one flow makes sense (the central flow)
  - We show empirically that this flow matches GD in a variety of DL settings

# Example: special case of 1 unstable eigenvalue

- We model the iterate as:

$$w_t = w(t) + x_t u_t$$

time-averaged  
iterate      top Hessian  
                  eigenvector

iterate      magnitude of  
                  oscillation

- Then the gradient is:

$$\nabla L(w_t) \approx \nabla L(w(t)) + x_t S(w(t)) u_t + \frac{1}{2} x^2 \nabla S(w(t))$$

gradient at time-  
averaged iterate      sharpness reduction

gradient at iterate      oscillation

- So the “time-averaged” gradient is:

$$\mathbb{E}[\nabla L(w_t)] \approx \nabla L(w(t)) + \mathbb{E}[x_t] S(w(t)) u_t + \frac{1}{2} \mathbb{E}[x^2] \nabla S(w(t))$$

gradient at time-  
averaged iterate

time-averaged gradient

variance of oscillations  
↓      sharpness reduction

oscillation

# Example: special case of 1 unstable eigenvalue

- We suppose that the time-averaged GD trajectory follows an ODE of the form:


$$\frac{dw}{dt} = -\eta \left[ \underset{\text{gradient flow}}{\nabla L(w)} + \frac{1}{2} \sigma^2(t) \underset{\text{sharpness penalty}}{\nabla S(w)} \right]$$

“instantaneous variance” of the oscillations  
(i.e. local time average of  $x^2$ )

- This flow averages out the oscillations, but keeps their *effect* on the trajectory.
- To determine  $\sigma^2(t)$ , we argue that only one value is possible.

- Empirically, the sharpness equilibrates at  $2/\eta$ .

- Therefore, we enforce that along the central flow,  $\frac{dS}{dt} = 0$ .



# Example: special case of 1 unstable eigenvalue

- The time derivative of the sharpness under our flow is:

$$\begin{aligned}\frac{dS}{dt} &= \left\langle \nabla S(w), \frac{dw}{dt} \right\rangle \quad \text{chain rule} \\ &= \left\langle \nabla S(w), -\eta \left[ \nabla L(w) + \frac{1}{2} \sigma^2(t) \nabla S(w) \right] \right\rangle \quad \text{substitute in our flow} \\ &= \left\langle \nabla S(w), -\eta \nabla L(w) \right\rangle - \frac{1}{2} \eta \sigma^2(t) \|\nabla S(w)\|^2 \quad \text{simplify}\end{aligned}$$

time derivative of sharpness  
under gradient flow

sharpness-reduction  
effect of oscillations

- We see that  $\frac{dS}{dt}$  is **affine** in  $\sigma^2(t)$ . In order for  $\frac{dS}{dt} = 0$ ,  $\sigma^2(t)$  must be:

$$\sigma^2(t) = \frac{2 \langle -\nabla L(w), \nabla S(w) \rangle}{\|\nabla S(w)\|^2}$$

# Central flow in action


- The central flow for a single unstable eigenvalue is:

$$\frac{dw}{dt} = -\eta \left[ \nabla L(w) + \frac{1}{2} \sigma^2(t) \nabla S(w) \right] \quad \text{where} \quad \sigma^2(t) = \frac{\langle -2 \nabla L(w), \nabla S(w) \rangle}{\| \nabla S(w) \|^2}$$

# Central flow in action

- The central flow for a single unstable eigenvalue is:


$$\frac{dw}{dt} = -\eta \left[ \nabla L(w) + \frac{1}{2} \sigma^2(t) \nabla S(w) \right] \quad \text{where} \quad \sigma^2(t) = \frac{\langle -2 \nabla L(w), \nabla S(w) \rangle}{\| \nabla S(w) \|^2}$$



# Central flow in action

- The central flow for a single unstable eigenvalue is:


$$\frac{dw}{dt} = -\eta \left[ \nabla L(w) + \frac{1}{2} \sigma^2(t) \nabla S(w) \right] \quad \text{where} \quad \sigma^2(t) = \frac{\langle -2 \nabla L(w), \nabla S(w) \rangle}{\| \nabla S(w) \|^2}$$



# Central flow in action

- The central flow for a single unstable eigenvalue is:


$$\frac{dw}{dt} = -\eta \left[ \nabla L(w) + \frac{1}{2} \sigma^2(t) \nabla S(w) \right] \quad \text{where} \quad \sigma^2(t) = \frac{\langle -2 \nabla L(w), \nabla S(w) \rangle}{\| \nabla S(w) \|^2}$$



# Central flow in action

- The central flow for a single unstable eigenvalue is:

$$\frac{dw}{dt} = -\eta \left[ \nabla L(w) + \frac{1}{2} \sigma^2(t) \nabla S(w) \right] \quad \text{where} \quad \sigma^2(t) = \frac{\langle -2 \nabla L(w), \nabla S(w) \rangle}{\| \nabla S(w) \|^2}$$



# Takeaways


- It's challenging to understand the oscillations in fine-grained detail
- But the *macroscopic* trajectory only depends on the *variance* of the oscillations
- This variance is easy to obtain
  - There is only one value that is compatible with the edge of stability equilibrium

# Interpretation as projection

- The central flow can be equivalently interpreted as a *projected gradient flow*:

$$\frac{dw}{dt} = -\eta \left[ I - \frac{\nabla S(w) \nabla S(w)^T}{\|\nabla S(w)\|^2} \right] \nabla L(w)$$

project out  $\nabla S(w)$  direction from  $\nabla L(w)$   
to keep sharpness  $S(w)$  fixed in place



# Complete central flow

- Similar to before, we make the ansatz that the time-averaged iterates follow:

$$\frac{dw}{dt} = -\eta \left[ \nabla L(w) + \frac{1}{2} \nabla_w \langle H(w), \Sigma(t) \rangle \right]$$

implicit curvature penalty



where  $\Sigma(t)$  models the  $\mathbb{E}[\delta_t \delta_t^T]$ , the covariance of the oscillations.

- We argue that only one value of  $\Sigma(t)$  is possible.



# Complete central flow

- We impose three conditions on the central flow:
  1. The flow should not increase any Hessian eigenvalues above  $2/\eta$
  2.  $\Sigma(t)$  should be supported within the Hessian's  $2/\eta$  eigenspace
  3.  $\Sigma(t)$  should be positive semidefinite
- These three conditions imply that  $\Sigma(t)$  must be the solution to a certain *cone complementarity problem*.
- The central flow is defined with this  $\Sigma(t)$ .


# Central flow in action




# Central flow in action




# Central flow can predict oscillation covariance




# Central flow can predict oscillation covariance




# Central flow can predict oscillation covariance




# Central flow can predict oscillation covariance




# Central flow can predict oscillation covariance



# Central flow can predict oscillation covariance




# Application: reasoning about loss curves




- The gradient descent loss curve is non-monotonic...

# Application: reasoning about loss curves




- The gradient descent loss curve is non-monotonic...
- ... but the *central flow* loss monotonically decreases:
$$\frac{dL(w(t))}{dt} \leq 0$$
- The central flow loss  $L(w(t))$  is a **potential function** for the optimization process.
- Its slope quantifies the speed of optimization.


# Application: reasoning about loss curves



- Loss is **higher** for GD than for central flow.
- Intuition: GD bounces between “valley walls”; central flow runs along “valley floor”



# Application: reasoning about loss curves



- The central flow models *both* the mean trajectory *and* the covariance of oscillations:


$$w_t \approx w(t) + \delta_t \quad \text{where} \quad \mathbb{E}[\delta_t] = 0, \quad \mathbb{E}[\delta_t \delta_t^T] = \Sigma(t)$$

- Thus, it can predict the *time-averaged* train loss of gradient descent:

$$\mathbb{E}[L(w_t)] \approx L(w(t)) + \frac{1}{\eta} \operatorname{tr} \Sigma(t)$$

time-averaged  
GD loss      loss along  
central flow      contribution  
from oscillations

# Application: reasoning about loss curves



- The central flow models *both* the mean trajectory *and* the covariance of oscillations:

$$w_t \approx w(t) + \delta_t \quad \text{where} \quad \mathbb{E}[\delta_t] = 0, \quad \mathbb{E}[\delta_t \delta_t^T] = \Sigma(t)$$

- Thus, it can predict the *time-averaged* train loss of gradient descent:


$$\mathbb{E}[L(w_t)] \approx L(w(t)) + \frac{1}{\eta} \text{tr} \Sigma(t)$$

time-averaged  
GD loss

loss along  
central flow

contribution  
from oscillations


# Application: reasoning about loss curves



- The central flow models *both* the mean trajectory *and* the covariance of oscillations:
$$w_t \approx w(t) + \delta_t \quad \text{where} \quad \mathbb{E}[\delta_t] = 0, \quad \mathbb{E}[\delta_t \delta_t^T] = \Sigma(t)$$
- Thus, it can predict the *time-averaged* train loss of gradient descent:
$$\mathbb{E}[L(w_t)] \approx L(w(t)) + \frac{1}{\eta} \text{tr} \Sigma(t)$$


time-averaged  
GD loss      loss along  
central flow      contribution  
from oscillations
- Both  $L(w(t))$  and  $\mathbb{E}[L(w_t)]$  are meaningful quantities to DL practitioners

# Central flow is the “true” training process




# A smooth curve is a simple object

- As a smooth curve, the central flow is a simple object.




# A smooth curve is a simple object

- As a smooth curve, the central flow is a simple object.
- The central flow update direction  $\frac{dw}{dt}$  reflects the near-term direction of motion.




# A smooth curve is a simple object

- As a smooth curve, the central flow is a simple object.
- The central flow update direction  $\frac{dw}{dt}$  reflects the near-term direction of motion.
- By contrast, the GD update  $-\eta \nabla L(w_t)$  is dominated by oscillations.





# Our analysis applies to generic neural nets



# Review

- Existing optimization theory does not apply in deep learning
  - Doesn't capture *cause and effect* for *deterministic gradient descent*
- But a different theory is possible
  - Deep learning objectives aren't that scary
  - Our analysis, while not rigorous, delivers accurate numerical predictions
  - Deep learning may call for a different approach than classical optimization



# What is the *goal* of optimization theory?

- Classically, a common goal is to characterize global rates of convergence.
  - But this might never be possible in deep learning
- Another goal is to characterize the local rate of convergence once near a minimum
  - But deep learning optimization doesn't occur near a minimum
- Our goal: characterize the *local dynamics throughout training*
  - These dynamics are (1) interesting, (2) important, and (3) generic.

# What is the *purpose* of an optimization paper?

- ML reviewers' favorite kind of paper: theoretical analysis + new SOTA algorithm
- But we are likely still in the theory-building stage
- *Basic* research now will enable SOTA algorithm design in the future

# **What *methods* are acceptable?**

- Optimization historically operates at a 100% level of mathematical rigor
- This standard may not be appropriate for deep learning
- People make assumptions that aren't true, so that they can leverage known proof techniques, rather than investigating what really happens
- The field should be comfortable with works at varying levels of rigor
- The right mathematical tools will develop gradually to fit the needs of the field

# A good field to work on

- Deep learning is one of the defining technologies of this century
- Optimization lies at the heart of deep learning
- There is room for an entire field on the theory of optimization in deep learning
- Applied mathematicians can help turn deep learning from alchemy to science

# Thanks to my collaborator Alex



Cohen\*, Damian\*, Talwalkar, Kolter, Lee. *Understanding Optimization in Deep Learning with Central Flows*. ICLR '25.

OpenReview:



Alex Damian

ArXiv: there's a draft on arXiv, but we're still putting the finishing touches on the final version

Email me for code: [jcohen@flatironinstitute.org](mailto:jcohen@flatironinstitute.org)

# PS: we also analyze Adam with $\beta_1 = 0$ (i.e. RMSProp)

- This algorithm doesn't make much sense according to traditional understandings, but works well in practice
  - How can we beat Adam if we don't understand it?
  - We show that understanding how Adam sets its dynamic preconditioner requires understanding its oscillatory EOS dynamics
  - We also show that Adam's efficacy relies on its ability to implicitly steer itself towards lower-curvature regions in which it can take larger steps
  - Part II of this talk: “How does Adam work?”
  - Thanks for listening!